

Comments on NIST AI 300-1 ipd

Guidance and Templates for Public-Facing AI Documentation — An AI Standards “Zero Draft,” Initial Public Draft (July 2026; NIST’s designation)

To: ai-standards+doczd@nist.gov

From: Jiajun Ma, independent researcher

Date: September 10, 2026

Capacity. Submitted in my personal capacity. The views are my own and represent no employer or institution.

Summary. One recommendation with replacement text on **Annex A fields 1.3.2 and 8.2.1**; a typed-absence convention; one editorial correction; brief responses on scoping, future topics, and process. Line numbers are the draft’s margin numbering.

Use of AI assistants. Per the Note to Reviewers: AI assistants (Anthropic Claude, Opus-class models, September 2026) were used to check quotations against the draft, search prior art, and edit prose. In the terms this comment proposes, the artifact is *automated-draft-human-edited* and was *checked-against-primary-sources*: I verified every quotation against the draft’s line-numbered text and every figure against the published data files; line numbers are given so it can be repeated.

Background. I maintain a public, DOI-registered corpus that classifies 84 institutional investors by stage of AI adoption from public evidence (doi.org/10.5281/zenodo.22247885; source and data at github.com/jima1835/ai-adoption-in-finance; live view at jima1835.github.io/ai-adoption-in-finance). A further 15 organizations were assessed and not classified: 10 for insufficient public evidence, 5 withdrawn after drafting, and are published by name. Each record is drafted by an AI research agent and verified by a person; the proposal, the decision, and every withdrawal are recorded in the artifact. Of the 84 records, 74 carry an agent proposal with a recorded outcome, 8 were human-originated, and 2 predate the provenance fields with no log entry. The corpus publishes agreement figures labelled as anchored (the reviewer saw the proposal first, so they are an upper bound); I do not rely on them. It is one deployment by one person, offered as a worked example, not as evidence that generalizes; it would itself fall under *automated-draft-human-edited*, and I have no commercial interest in the outcome.

1. The gap

Clause 1 (lines 110–113) states the templates are “requirements for documentation artifacts in the form of templates against which conformity of documentation artifacts can be assessed.” Clause 5.1 (lines 598–599) defines that conformity: “All fields designated in the templates as required shall be present and populated.” The only Required root field in either template is *Identifying Descriptors*, whose subfields concern identity and version; no required field in either template or default profile concerns how the artifact itself was produced or checked. It follows that an artifact generated by an AI system and approved without substantive review is fully conformant, and nothing the template requires would distinguish it from one written by domain experts.

Two fields are where the difference would show, and neither can show it. Annex A field 1.3.2, *Documentation Authors* — “One or more names of entities who prepared the documentation artifact” — is Optional in both default profiles, and a name cannot separate drafting from verification, or human from automated production. Field 8.2.1, *Documentation Quality Control* (Recommended), “May include a record of when and by whom the accuracy of the documentation was reviewed.” It never asks against *what* — the underlying evidence, or the generated text — though the pre-annotation and automation-bias literatures (Fort and Sagot 2010; Lingren et al. 2014; Skitka, Mosier and Burdick 1999) show that condition changes what a review is worth.

The templates this draft’s pattern descends from do not close the gap either. Model Cards (Mitchell et al. 2019) and Datasheets for Datasets (Gebru et al. 2021) document the model or dataset; neither has a field for how the card itself was produced or checked, and Hugging Face’s model-card template carries only “Model Card Authors” — names, with the same limitation as 1.3.2. I have not surveyed attestation standards outside that lineage, nor the comments on the September 2025 extended outline. The Note to Reviewers asks commenters to disclose AI assistance; the templates do not ask it of documentation artifacts.

2. Recommendation and replacement text

No change to Clause 5. Both additions are Annex A subfields, which Clause 5.1 (lines 605–609) says exist because the default profile “can accommodate more prescriptiveness.”

(a) Add subfield 1.3.2.1 *Production Method*, designated Recommended (1.3.2 currently has no subfields, so no renumbering follows), with a closed value set:

- `human-authored` — written by people, with or without non-generative tooling
- `automated-draft` — generated by an automated system and released without substantive human editing
- `automated-draft-human-edited` — generated by an automated system and substantively edited by people
- `not-recoverable` — the artifact predates the provider’s recording of production method

Elevate 1.3.2 from Optional to Recommended so the subfield has a populated parent (Clause 6, lines 662–664: a subfield’s parent shall be designated at least as strongly under the same conditions).

(b) Add subfield 8.2.1.1 *Verification Condition*, designated Recommended (8.2.1 has no subfields), with a closed value set:

- `checked-against-primary-sources` — a reviewer compared the artifact’s claims to underlying evidence (data, code, evaluation outputs) independently of the artifact’s text
- `checked-against-generated-text` — a reviewer read the artifact as drafted and accepted, revised, or rejected it
- `not-independently-checked`
- `not-recoverable`

Design notes. The two axes are deliberately separate: who *produced* the content and what the *checker* could see are different facts, and one “human-reviewed” indicator conflates them. The value sets are deliberately small: a contributor-role taxonomy on the model of CRediT (Brand et al. 2015) is the natural extension, but the experience in §3 suggests that an enumeration specified ahead of practice under-specifies, and that a few values which can be populated honestly are worth more than many that cannot. The specific value sets are a starting proposal for NIST to test against providers. Both are artifact-level; either may be repeated per root field where production or checking differed. Recommended, not Required, is the deliberate ask: the draft keeps requirements minimal and lets conformity schemes require what a profile recommends (Clause 6), so the gap closes for providers who populate the field and becomes visible for those who do not. The values are unaudited self-report, as every Annex A field is; a false value is a false statement in a public artifact, which is a different position from silence, and 8.1.1 *Review Processes* is where an audit profile could later attach. Preferred combination: adopt the general absence convention in §3 and drop the `not-recoverable` value from both sets; if the convention is not adopted, keep the field-level value, since these two fields are where a legacy cohort is certain.

On scope. Clause 5.1 (lines 612–613) says “the templates cannot capture organizational process elements,” but the draft already treats *descriptions* of process as artifact content — 8.1.1, 8.2 (“the organizational practices that yielded the model documentation”), and 8.2.1 itself. This adds one more of the same kind.

Editorial. The table of contents (lines 29–31) orders 4.2.1 comprehensibility, 4.2.2 informativeness, 4.2.3 correctness; the body (lines 385, 392, 400) orders 4.2.1 correctness, 4.2.2 comprehensibility, 4.2.3 informativeness. Recommend the body be renumbered to match the table of contents, and cross-references checked.

3. Why this draft cycle, not a later profile

A provenance field added after artifacts exist can be populated only where a contemporaneous log happens to exist; where none does, it never can be. In the corpus above, a four-value provenance vocabulary (`agent-proposed-accepted`, `human-revised`, `human-originated`, `unknown-pre-capture`) was introduced after review had begun. For 20 of the 74 records with a recorded proposal, the value could be reconstructed only because a review tool had kept a decision log; for 2 of the 84 records no entry exists and the value is permanently unrecoverable. One deployment, one domain, one reviewer: I claim no rate from it, only the mechanism.

The same vocabulary failed in two documented ways that shaped §2: it could not record that the reviewer, not the agent, had found the decisive source for a row, nor that the reviewer had directed a row’s inclusion, because the value describes only the label. Both argue for two axes rather than one provenance token.

Precedent. Two regulatory precedents show different things. FDA’s 1997 electronic-records rule (21 CFR Part 11) was narrowed by 2003 guidance granting enforcement discretion for “systems that were operational before August 20, 1997”: relief for a legacy cohort is possible, but only from an empowered regulator, after the fact, at a cost. The EU AI Act’s documentation obligations for general-purpose models apply from August 2, 2025 (Art. 113(b)), with models already on the market given until

August 2, 2027 (Art. 111(3)): a legacy cohort exists in that regime now. A voluntary template has no authority that can grant relief later. Its one move is to specify the value at the outset, and a not-recoverable value is what lets a legacy model conform honestly rather than not at all.

Typed absence. The draft itself creates several reasons a field may be empty and no way to record which: information withheld under 4.2.4 (lines 412–414, “artifacts should omit private, proprietary, or legally impermissible information”); information “not available to or cannot be revealed by the model provider” (A.3 preamble, lines 733–736); an optional field simply not populated; and no provision for artifacts predating the provider’s process. Two dataset fields already permit “unknown” ad hoc (Table A.2.1, 4.3.3 and 4.4). I recommend one convention, inserted in Clause 5.1 after the paragraph on field population (lines 600–603) and applicable to any Annex A field: “Where a field is present but not populated with substantive content, it shall state the reason from: not-applicable / not-assessed / withheld / not-recoverable.” Applied to Table A.2.1 field 4.3.3, the guidance “May be indicated as approximate or unknown” would read “May be indicated as approximate; absence shall be typed per Clause 5.1.” Two precedents do different work. HL7’s NullFlavor code system, carried into FHIR as data-absent-reason, shows a voluntary technical standard can type absence at the field level. For not-applicable and not-assessed, FDA’s labeling rule (21 CFR 201.57) is the precedent: the Contraindications section must state “None” when there are none, because silence is not an admissible answer; for not-recoverable, the legacy-cohort precedents above are. A 2024 analysis of 32,111 Hugging Face model cards (Liang et al.) found the limitations and evaluation sections least filled — optional sections, which is what Optional designations create. A template cannot force them full; it can make an empty one say why.

Why these fields are likelier to be populated than other optional fields. Three reasons, offered as reasons rather than a prediction. First, the provider that documents well has no way, under the current draft, to show that it did; a production-method field is the one entry that separates an expert-written artifact from a generated one, so such providers have an interest in populating it that they do not have in a free-text limitations section. Second, each field is one value from a closed set, and a typed absence costs less than the substantive field it stands in for; the sections Liang et al. found empty are free text. Third, beyond those, uptake depends on an adopting party conditioning on the field (§4), and I do not claim otherwise.

4. Responses to the pilot project’s other questions

Scoping. Annex C, “Relationships with Other ISO/IEC Standards on AI,” is under consideration. The EU’s mandatory documentation requirements for general-purpose models (Art. 53; Annex XI; the Commission’s July 24, 2025 training-content template) should be crosswalked there; a voluntary template neither a superset of nor mapped to a mandatory one will not be adopted by providers serving both regimes. §2 is distinct from content-provenance work such as C2PA and NIST AI 100-4, which mark the provenance of *content* a system produces, not of the *documentation* that describes it.

Future topics. Clause 1 (lines 119–122) excludes system-level documentation as “less mature.” My experience corroborates the immaturity and suggests it is the reason to standardize, not to wait: classifying 84 organizations from public disclosure required 276 dated evidence items from 116 source domains, and the 10 organizations that could not be classified fell into a category — insufficient public evidence — that cannot separate “no adoption” from “no disclosure.” Deployer-side documentation — what an organization adopting a third-party model does with it, under what oversight, at what stage — is where that cost sits, and a template would create the practice the draft says does not yet exist.

Process. Three frictions from a first-time, unaffiliated commenter: the “optional commenting template” the Note to Reviewers encourages is not linked from the publication or pilot-project page; NIST’s concurrent AI drafts route to three addresses and the pilot page’s “TEVV draft” is a different document from NIST AI 200-2; and this draft and AI 200-2 share subject matter and an author but close three weeks apart, so cross-cutting comments must be split. On adoption: trial registration rose sharply once journals made it a condition of publication (ICMJE 2005), while statutory results reporting under FDAAA 801 remained substantially incomplete a decade later (Anderson et al. 2015; DeVito et al. 2020). The obligations differ, but the pattern suggests templates get populated when a party the provider depends on conditions on them. NIST is not that party; the final draft might name who it expects to be — a model hub, a procurement rule, a publication venue.

Availability. The corpus, agreement file (data/agreement.json), negative record (data/not_classified.json), and transition log (data/transitions.jsonl) are public (CC BY 4.0 data, MIT code) at github.com/jima1835/ai-adoption-in-finance, archived at doi.org/10.5281/zenodo.22247885, and may be used freely as a worked example — including of how a provenance vocabulary under-specifies.

References

- Anderson, M. L., et al. (2015). Compliance with results reporting at ClinicalTrials.gov. *N. Engl. J. Med.* 372(11): 1031–1039.
- Brand, A., Allen, L., Altman, M., Hlava, M., and Scott, J. (2015). Beyond authorship: attribution, contribution, collaboration, and credit (the CRediT taxonomy). *Learned Publishing* 28(2): 151–155.
- Geburu, T., et al. (2021). Datasheets for datasets. *Commun. ACM* 64(12): 86–92.
- DeVito, N. J., Bacon, S., and Goldacre, B. (2020). Compliance with legal requirement to report clinical trial results on ClinicalTrials.gov. *Lancet* 395(10221): 361–369.
- FDA (2003). Part 11, Electronic Records; Electronic Signatures — Scope and Application (guidance). 68 FR 52779.
- Fort, K., and Sagot, B. (2010). Influence of pre-annotation on POS-tagged corpus development. *Proc. Fourth Linguistic Annotation Workshop (ACL)*, 56–63.
- HL7 International. NullFlavor code system (v3); FHIR data-absent-reason. terminology.hl7.org/CodeSystem-v3-NullFlavor.html
- ICMJE (2005). Update on trials registration. https://www.icmje.org/news-and-editorials/update_2005.html
- Liang, W., et al. (2024). What's documented in AI? Systematic analysis of 32K AI model cards. *arXiv:2402.05160*.
- Lingren, T., et al. (2014). Evaluating the impact of pre-annotation on annotation speed and potential bias. *J. Am. Med. Inform. Assoc.* 21(3): 406–413.
- Mitchell, M., et al. (2019). Model cards for model reporting. *Proc. FAT* 2019*, 220–229.
- Ma, J. (2026). AI Adoption in Finance: a human-verified, stage-classified corpus and its human–AI disagreement record. github.com/jima1835/ai-adoption-in-finance; Zenodo doi:10.5281/zenodo.22247885
- Regulation (EU) 2024/1689 (AI Act), Arts. 53, 111(3), 113(b); Annex XI; Commission template for public summaries of training content (July 24, 2025). 21 CFR § 201.57; 21 CFR Part 11.
- Skitka, L. J., Mosier, K. L., and Burdick, M. (1999). Does automation bias decision-making? *Int. J. Human-Computer Studies* 51(5): 991–1006.