

Comments on NIST AI 200-2 ipd — The TEVV-Athlon Framework for Evaluating AI Systems (Initial Public Draft, August 2026)

To: TEVV-Athlon@nist.gov · Date: 19 September 2026

Respondent: Jiajun Ma, independent researcher

Capacity: Submitted in my personal capacity; the views are mine and represent no employer or institution. The letter contains no proprietary information and may be released in full.

Seat: The front matter invites comment from “users of AI evaluation reports – for instance, business decision-makers, procurement specialists, researchers, and technical staff.” I write as one of those, from the reader’s side. I build and maintain a public dataset (<https://doi.org/10.5281/zenodo.22247885>) that places 84 institutional investors on a four-stage AI-adoption scale using public evidence only; an AI research agent drafts each record, a person verifies it, and the record of where the two disagreed is published with the data. The pipeline enforces a recusal list: institutions with which I have an affiliation are never assessed. A paper on the verification design has been accepted at the AAAI 2026 Fall Symposium on AI Trustworthiness and Risk Assessment for Challenged Contexts (ATRACC). I also commented on NIST AI 300-1 ipd on 10 September 2026 (<https://doi.org/10.5281/zenodo.22701607>); the two letters address different documents and make different requests.

Summary

The framework’s central choice, that an organization builds its own evaluation from stated objectives rather than adopting a generic benchmark, is right. Its cost is that two TEVV-Athlons run by two organizations will not be readable side by side unless something is held fixed. I recommend fixing the report, not the evaluation: one primary recommendation with replacement text for §2.4, one secondary recommendation with replacement text for §2.1, one note, and a statement of what my own evidence does and does not show. Nothing proposed prescribes which Blocks, Events, or Tools an organization chooses.

1. Primary: §2.4 should specify what a TEVV-Athlon reports to people who did not run it

What the draft says. §2.4 (lines 305–315) describes synthesis as reviewing the stages in reverse order to “generate relevant insights for the organization,” and states that the synthesis “influences organizational decisions regarding the assessed AI system.” It names no reader, no minimum content, and no format. §2.1 asks, in the second of its seven boxed questions (between lines 240 and 241), “Who will care about the results of this assessment?” Nothing later requires the output to be legible to whoever the answer is. Appendix D, Table 6, lists “Reporting limitations” and “Measurement of uncertainty” among considerations that are “not intended as prescriptive requirements” (lines 776–777). The draft’s own example is adapted from NIST’s ARIA 0.1 pilot (footnote 1, marked at line 323; NIST AI 700-2). That pilot report names the sessions it excluded and why, and gives its denominators [7, pp. 1–2, 4, 6, 22]; it does not state what annotators were shown before labeling, reports disagreement counts but no agreement statistic, gives no system version or date, and states no condition for repeating the assessment [7, pp. 6, 10, 13, 22]. Part of what I ask for below is already NIST practice; the framework does not carry it forward as a reporting expectation.

Why it matters. Most readers of an evaluation did not run it and cannot re-run it: boards, investment committees, supervisors, and allocators assessing a manager or a vendor from what it publishes. The FSB’s June 2026 consultation observes that third-party providers “may change the underlying models powering their AI systems without adequate notice, leaving financial institutions unaware of changes” [1]. On 10 September 2026 the SEC Chairman told the Commission’s Investor Advisory Committee that “widespread reliance on similar tools can allow errors to cascade across the marketplace,” and that the Commission “will not prescribe the specific models that firms must employ—we are not a merit regulator” [2]. If the regulator will not rate the models, the evaluation report is what the market has, and a report that cannot be read across organizations does not do that job. The problem is not specific to finance: the draft’s own example produces a Violation Frequency for a chatbot that will be read by a product owner who did not run the red team, and a June 2026 proposal for “evaluation cards” reaches the same conclusion for benchmark reporting in general [8].

Recommendation. Insert a subsection 2.4.1 after line 315:

2.4.1 Reporting. A TEVV-Athlon whose results will be relied on by readers who did not design or conduct it should produce a summary for those readers. At minimum the summary states: (a) the assessment objectives and the AI system’s declared stage from Stage 1; (b) the Blocks measured, and any Blocks considered and not measured, each with the reason; (c) for each Event, the Tool used, the source of the data, and, where a person’s judgment produced the data, whether the person judged from the raw record alone or was shown a prior label, score, or proposed answer before judging, together with the number of items judged and the number excluded and why; (d) the version of the AI system and of the TEVV-Athlon, the date, and the condition under which the assessment is to be repeated; (e) known limitations, in the evaluators’ own words; (f) who conducted the assessment, by role or function, and whether they served as front-line developers of the system or are independent of the team that developed or deployed it. Without (b), (c) and (f), a reader outside the evaluation cannot tell what a reported statistic covers, what the judge was shown, or whose judgment it was.

And add one row to Table 6: *Reporting for non-evaluators* — Provide a summary stating objectives, declared stage, Blocks measured and not measured, per-Event Tool, data source and judgment condition, version and repeat condition, limitations, and the evaluators’ relationship to the system’s developer.

The elements are not untested. When a credit rating agency takes a rating action, SEC Rule 17g-7 requires a form stating the rating’s potential limitations, “the reliability, accuracy, and quality of the data relied on,” any “limits on accessibility to certain documents or other types of information that would have better informed the credit rating,” who paid for the rating, and a signed attestation that it “was an independent evaluation” [11]. That rule is mandatory and addresses a paid outsider rating an issuer, while 2.4.1 is voluntary and its concern is an organization assessing its own system; the precedent shows only that these report elements can be specified and have been, for an evaluation consumed by readers who did not produce it.

Three of the six items consolidate content the draft already asks for somewhere: (a) restates Stage 1 (and, if the §2.1 change below is not adopted, reads as the lifecycle stage the draft already asks evaluators to consider), (d) extends Appendix E’s “Documentation of procedures,” and (e) restates Table 6’s “Reporting limitations.” Items (b), (c) and (f) are new to this draft. On (b): the AI RMF already expects that “the risks or trustworthiness characteristics that will not – or cannot – be measured are properly documented” (MEASURE 1.1) [9]; the draft says its final stage feeds the RMF’s Manage function (line 315) but never asks the report to carry that documentation. On (c): the draft’s example measures Violation Frequency through “human annotators according to an Annotation Schema” (lines 393–394) and reports “descriptive statistics” (lines 399–402) without stating whether annotators labeled from the transcript alone or from a pre-filled label, and the pilot report it draws on is silent on the same point [7, p. 6]. Whether a judge worked from the raw record or from a proposed answer changes what the resulting number means; the effect is documented for automation bias and for pre-annotated labeling [3–5]. On (f): the AI RMF expects “internal experts who did not serve as front-line developers for the system and/or independent assessors” to take part in assessments (MEASURE 1.3) [9], and Table 6 lists “Independent review or challenge” as a consideration; neither asks the report to say who conducted the TEVV-Athlon. Stating the relationship is a floor, not a test: the banking agencies’ revised model risk guidance holds that effective challenge needs “sufficient independence to maintain objectivity” and that accountability covers “potential conflicts of interest (e.g., misalignment of incentives between different reporting lines, such as model development and validation groups),” while also placing the quality of a review in its rigor “rather than on organizational structure” [6]. A reader can weigh rigor only if the report first says who did the work; and that guidance, as noted above, does not itself cover generative or agentic systems.

2. Secondary: §2.1 should require a declared stage on a published scale, with design proportional to it

What the draft says. Lines 231–233: “evaluators should consider where the AI system is in its lifecycle. The lifecycle stage may influence the TEVV-Athlon design, as it can affect the feasibility of certain assessment methods and the intended use of assessment results.” Fig. 1 (page 2, caption between lines 178 and 179) supplies a lifecycle classification “based on an OECD taxonomy,” which is a position in the technical lifecycle. §4.2 states the proportionality principle (lines 440–442): “a risk-based approach to measurement can be applied: systems with higher potential risk need more rigorous measurement.” Nothing asks the organization to declare the risk or reliance level it judged, on what scale, or to report it. A reader of the result cannot check that the measurement was proportionate to anything.

Why it matters. Depth of reliance is a different variable from lifecycle position. The same deployed model warrants a different assessment inside a contained pilot than inside decisions the organization cannot easily unwind. Financial

supervisors have asked for proportionality without supplying the scale: the FSB’s Sound Practice 9 asks institutions to evaluate AI performance “proportionately to their materiality and risk” [1], and the US banking agencies’ April 2026 model risk guidance aligns “the nature and rigor of validation” with “the model’s approach, use, and materiality,” while placing generative and agentic AI outside its scope [6]. A sector framework built with NIST input already does what I ask: the Cyber Risk Institute’s Financial Services AI Risk Management Framework (February 2026, developed with the Financial Services Sector Coordinating Council and released through the US Treasury) “begins by identifying an organization’s AI adoption stage” on a four-stage scale and maps its 230 control objectives to the declared stage [10]. My dataset is the external-observation counterpart, the same staged form applied from public disclosure alone with published boundary rules. NIST need not adopt either scale. The requirement is that a scale be declared and reported; the declaration is a self-report, and reporting it under 2.4.1(a) with the scale’s boundary rules is what lets a reader contest it.

Recommendation. Replace lines 231–233 with:

Evaluators should state the AI system’s stage of adoption or organizational reliance on a published ordinal scale with written boundary rules, and record it as an input to the TEVV-Athlon. The number and depth of Blocks, Events, and Tools selected in Stages 2 and 3 should be proportionate to the declared stage, consistent with §4.2. A change in the declared stage is a condition for repeating the assessment (see 2.4.1(d)).

3. One note: the cross-reference to AI 300-1, and why 2.4.1 belongs here

A 2.4.1 summary and a model’s public documentation under NIST AI 300-1 (Clause 5, Table 2, root field 6, Evaluation) describe the same evaluation from two sides; one sentence in each draft pointing to the other would let a reader place them side by side. They are not substitutes: AI 300-1 documents a model for the public, while a TEVV-Athlon is built by one organization for its own objectives and read by boards, committees, and supervisors who will never see a model card. §2.4 is the only point at which the framework hands results to any reader, which is why the minimum belongs there.

4. What my evidence shows, and what it does not

The dataset behind this comment is one deployment in one domain with one reviewer, who is also its producer and the author of this letter, and it is self-published. No rate is claimed from it. On item (c): 74 of the 84 records came with an agent proposal and a recorded decision; 8 were written by the reviewer without one; 2 predate the provenance record. On those 74 the reviewer saw the agent’s proposed classification and rationale before deciding, and the published record labels the resulting agreement as anchored, meaning the reviewer had seen a proposed answer first. How anchored agreement and independent agreement differ, in size or in direction, was not measured; a blind re-code by an independent coder is scheduled for the fourth quarter of 2026. In my 10 September comment on AI 300-1 I described these figures as an upper bound. That asserted a direction that was not measured; the dataset’s release of 19 September 2026 withdraws the wording, and I withdraw it here. On item (b): the dataset publishes, by name and with reasons, the entities assessed under the same rules and not classified, because the classified rows cannot be read without them. Each practice is one published file. Neither is asked for in the draft, and NIST’s own pilot report does the second and not the first.

Disclosure. AI assistants (Anthropic Claude, September 2026) were used to draft and check this letter, and three AI-generated adversarial reviews were run on each draft, including this one. I verified every quotation against the three NIST PDFs cited at the line and page numbers given, once, myself; the quotations from the FSB, SEC, banking-agency and SEC-rule sources were checked against the primary texts at the URLs given. The numbers are supplied so the check can be repeated.

Jiajun Ma

Independent researcher

jiajunma1835@gmail.com

<https://jima1835.github.io/ai-adoption-in-finance> · <https://github.com/jima1835/ai-adoption-in-finance>

References

- [1] Financial Stability Board, *Sound Practices for Responsible Adoption of Artificial Intelligence (AI): Consultation report*, 10 June 2026, <https://www.fsb.org/uploads/P100626.pdf> (subsection “Third-party dependencies and concentration”; Sound Practice 9, “Performance management”).
- [2] P. S. Atkins, *Remarks at the Investor Advisory Committee Meeting*, U.S. Securities and Exchange Commission, 10 September 2026, <https://www.sec.gov/newsroom/speeches-statements/atkins-remarks-iac-091026>.
- [3] L. J. Skitka, K. L. Mosier, and M. Burdick, “Does automation bias decision-making?”, *International Journal of Human-Computer Studies* 51(5), 1999, 991–1006.
- [4] Z. Buçinca, M. B. Malaya, and K. Z. Gajos, “To trust or to think: Cognitive forcing functions can reduce overreliance on AI in AI-assisted decision-making,” *Proc. ACM Hum.-Comput. Interact.* 5(CSCW1), 2021, Art. 188.
- [5] K. Fort and B. Sagot, “Influence of pre-annotation on POS-tagged corpus development,” *Proc. Fourth Linguistic Annotation Workshop (LAW IV)*, ACL, 2010.
- [6] Board of Governors of the Federal Reserve System, Office of the Comptroller of the Currency, and Federal Deposit Insurance Corporation, *Supervisory Guidance on Model Risk Management* (revised), 17 April 2026, issued as SR 26-2 and OCC Bulletin 2026-13, <https://www.federalreserve.gov/supervisionreg/srletters/SR2602.pdf> (Section III, effective challenge; Section V, validation; Section VI, roles and responsibilities; footnote 3, generative and agentic AI).
- [7] NIST, *Assessing Risks and Impacts of AI (ARIA): ARIA 0.1 Pilot Evaluation Report*, NIST AI 700-2, November 2025, <https://doi.org/10.6028/NIST.AI.700-2>.
- [8] EvalEval Coalition, “Evaluation Cards: An Interpretive Layer for AI Evaluation Reporting,” arXiv:2606.09809, June 2026, <https://arxiv.org/abs/2606.09809>.
- [9] NIST, *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*, NIST AI 100-1, January 2023, MEASURE 1.1 and MEASURE 1.3, <https://doi.org/10.6028/NIST.AI.100-1>.
- [10] Cyber Risk Institute, *Financial Services AI Risk Management Framework*, developed with the Financial Services Sector Coordinating Council and released through the U.S. Department of the Treasury, February 2026, <https://cyberriskinstitute.org/artificial-intelligence-risk-management/> (AI Adoption Stage Questionnaire; Risk and Control Matrix “organized by adoption stage”).
- [11] 17 CFR § 240.17g-7(a)(1)(ii)(D), (E), (J) and (a)(1)(iii), *Disclosure requirements* (nationally recognized statistical rating organizations), <https://www.ecfr.gov/current/title-17/section-240.17g-7>.